

# Sarvam AI

**Assignment:** Building a RAG system and an Agent that can perform smart actions based on the user's query

## RAG -

A **Retrieval-Augmented Generation (RAG) Agent** is a hybrid model that combines two powerful methods for answering user queries:

1. **Retrieval:** The agent retrieves relevant information from a knowledge base, document set, or vector database.
2. **Generation:** The agent uses a language model ( Cohere) to generate a response based on the retrieved information.

A RAG agent aims to improve the quality of answers by combining factual information retrieval and generative capabilities, ensuring that the model's responses are accurate, relevant, and grounded in external knowledge.

## How does a RAG work?

1. **Input Query:** The user submits a question.
2. **Retrieval Step:** The agent searches for relevant chunks of text from a document repository or vector database that could help answer the question. These chunks are retrieved based on similarity with the query.
3. **Generation Step:** Once relevant documents or text segments are retrieved, a language model generates a response based on the content of the retrieved information.
4. **Answer:** Then it returns the generated answer, which is more accurate because it is grounded in specific retrieved information rather than purely based on the model's internal knowledge.

## Components of an RAG with Agent

1. **Retriever:**
  - Based on the user query, retrieve relevant documents or text chunks from a corpus or knowledge base.
  - Often uses vector databases like **Chroma** that store document embeddings to make retrieval efficient and precise.
2. **Language Model (LLM):**
  - A pre-trained model (e.g., Cohere, GPT) that can generate natural language answers.
  - The retrieved documents are passed to the language model, which processes the information and generates a coherent answer.
3. **QA Chain:**

- Combines the retrieval and generation steps.
  - The retrieved documents are used as input for the language model to generate responses in a contextually grounded way.
4. **Tool & Agent:**
- A **tool** defines the functionality of the RAG chain.
  - An **agent** is a higher-level interface that can use multiple tools and determine which one to use for a given task.

## Walkthrough of the Code

Here's an explanation of how the provided code creates and utilizes an RAG and agent-

- **Loading Documents:** The first step is to load a document (in this case, a PDF) into memory. This document serves as the knowledge base from which relevant information will be retrieved.

```
loader = PyPDFLoader("/content/iesc111.pdf")
docs = loader.load()
```

- **Splitting Documents:** The document is split into smaller chunks of text because trying to retrieve large documents as a whole might lead to inefficient or inaccurate retrieval.

```
text_splitter = CharacterTextSplitter(chunk_size=800,
chunk_overlap=50)

docs_split = text_splitter.split_documents(docs)
```

- **Creating Embeddings:** Next, the code uses **Cohere Embeddings** to convert these document chunks into vector representations. This allows the agent to compare the user's query with the document chunks efficiently.

```
embeddings = CohereEmbeddings(model="embed-english-v3.0")
```

- **Storing in a Vector Database:** The document embeddings are stored in **Chroma**, a vector database that enables retrieval based on the user's query.

```
vectoradb = Chroma.from_documents(documents=docs_split,
embedding=embeddings, persist_directory=persist_directory)
```

- **Setting up the Retriever:** A retriever is set up to handle querying the vector database. It will pull relevant document chunks based on how similar they are to the input query.

```
retriever = vectordb.as_retriever()
```

- **LLM and RetrievalQA Chain:** A **Retrieval-Augmented Generation (RAG)** chain is created using the Cohere language model and the retriever. When the agent needs to answer a question, it will first retrieve relevant chunks of text and then use the language model to generate an answer.

```
llm = ChatCohere()
rag_chain = RetrievalQA.from_chain_type(
    llm=llm,
    chain_type="stuff", # "stuff" means using retrieved
documents directly
    retriever=retriever
)
```

- **Creating a Tool:** A **tool** is defined for the agent. The tool encapsulates the functionality of the RAG chain, making it available for the agent to use.

```
rag_tool = Tool(
    name="DocumentQA",
    func=rag_chain.run,
    description="Use this tool to answer questions by
retrieving and using context from documents."
)
```

- **Initializing the Agent:** The **agent** is initialized with the RAG tool. It is a **Zero-Shot React** agent, which means it decides which tool to use based on the description of the task.

```
agent = initialize_agent(
    tools=[rag_tool],
    agent=AgentType.ZERO_SHOT_REACT_DESCRIPTION,
    llm=llm,
    verbose=True
)
```

- **Querying the Agent:** Finally, the code allows the user to input a query, which is processed by the agent. The agent retrieves relevant documents from the vector database, generates an answer using the language model, and returns it to the user.

```
def query_rag_chain(query):  
    response = agent.run(query)  
    print(response)
```

